

dqvsd9zbp

April 26, 2024

1 Importing required libraries

```
[97]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter
import math
from scipy import sparse
import warnings
from sklearn.metrics.pairwise import cosine_similarity
from surprise import Reader, Dataset
warnings.filterwarnings("ignore")
```

2 Importing dataset

We have 4 files in total in the netflix directory, we will now import these 4 files into our code and combine them into a single file

```
[ ]: full_data = open('data.csv', mode='w')
files = ['netflix/combined_data_1.txt', 'netflix/combined_data_2.txt', 'netflix/
combined_data_3.txt', 'netflix/combined_data_4.txt']
for file in files:
    with open(file) as f:
        for line in f:
            line = line.strip()
            if line.endswith(':'):
                movie_id = line.replace(':', '')
            else:
                row = [x for x in line.split(',')]
                row.insert(0, movie_id)
                full_data.write(','.join(row))
                full_data.write('\n')
full_data.close()
```

- We stored the combined data in data.csv file in our project folder
- Now we will create a pandas data frame from our csv file

```
[98]: full_data = pd.read_csv('data.
↳csv',sep=',',names=['movie_id','user_id','rating','date'])
full_data.head()
```

```
[98]:   movie_id  user_id  rating      date
0         1  1488844        3  2005-09-06
1         1   822109        5  2005-05-13
2         1   885013        4  2005-10-19
3         1    30878        4  2005-12-26
4         1   823519        3  2004-05-03
```

We will now convert the date column to date type because it is of string type now

```
[99]: full_data.date = pd.to_datetime(full_data.date,format='%Y-%m-%d')
full_data.sort_values(by='date',inplace=True)
```

```
[100]: full_data.describe()
```

```
[100]:
```

	movie_id	user_id	rating	date
count	3.593923e+07	3.593923e+07	3.593923e+07	35939231
mean	3.359009e+03	1.322287e+06	3.595750e+00	2004-10-12 07:56:38.656124928
min	1.000000e+00	6.000000e+00	1.000000e+00	1999-11-11 00:00:00
25%	1.798000e+03	6.608560e+05	3.000000e+00	2004-05-03 00:00:00
50%	3.427000e+03	1.318672e+06	4.000000e+00	2005-01-20 00:00:00
75%	4.996000e+03	1.984358e+06	4.000000e+00	2005-07-05 00:00:00
max	6.386000e+03	2.649429e+06	5.000000e+00	2005-12-31 00:00:00
std	1.863574e+03	7.645951e+05	1.085413e+00	NaN

There are no duplicates/NaN values in the dataset as provided by the netflix

Basic Statistics

```
[101]: print("Total no of ratings :",full_data.shape[0])
print("Total No of Users   :", len(np.unique(full_data.user_id)))
print("Total No of movies  :", len(np.unique(full_data.movie_id)))
```

Total no of ratings : 35939231

Total No of Users : 476041

Total No of movies : 6386

Splitting the Data into train and test splits

```
[102]: train_data = full_data[full_data['user_id'] < 10000]
```

```
[103]: train_data.shape
```

```
[103]: (130782, 4)
```

```
[104]: test_data = full_data[(full_data['user_id'] > 10000) & (full_data['user_id'] < 13000)]
```

```
[105]: test_data.shape
```

```
[105]: (43448, 4)
```

Basic train data statistics

```
[106]: print("Total no of ratings :", train_data.shape[0])
print("Total No of Users      :", len(np.unique(train_data.user_id)))
print("Total No of movies     :", len(np.unique(train_data.movie_id)))
```

Total no of ratings : 130782

Total No of Users : 1803

Total No of movies : 4871

Basic test data statistics

```
[107]: print("Total no of ratings :", test_data.shape[0])
print("Total No of Users      :", len(np.unique(test_data.user_id)))
print("Total No of movies     :", len(np.unique(test_data.movie_id)))
```

Total no of ratings : 43448

Total No of Users : 517

Total No of movies : 3449


3 Exploratory Data Analysis

```
[108]: def beautify(num, units = 'M'):
units = units.lower()
num = float(num)
if units == 'k':
    return str(num/10**3) + " K"
elif units == 'm':
    return str(num/10**6) + " M"
elif units == 'b':
    return str(num/10**9) + " B"
```

3.1 Histogram of ratings

```
[109]: counter = Counter(full_data.rating)
items, counts = zip(*counter.items())
plt.bar(items, counts)
plt.xlabel('Items')
plt.ylabel('Counts (in millions)')
max_count_millions = math.ceil(max(counts) / 1_000_000)
counts_millions = [count / 1_000_000 for count in counts]
```

```
plt.yticks([count * 1_000_000 for count in range(0, max_count_millions + 1)],  
           ↳ ['{}M'.format(count) for count in range(0, max_count_millions + 1)])  
  
plt.show()
```



```
[110]: counter
```


```
[110]: Counter({4: 12080581, 3: 10356248, 5: 8167588, 2: 3664652, 1: 1670162})
```

We can see that the no.3 and no.4 ratings are more in numbers. We can easily predict that the average rating of all movies will be in between 3 and 4

3.2 No of ratings per month from the beginning of time in dataset to the end

```
[111]: train_data['date'] = pd.to_datetime(train_data['date'])  
train_data.set_index('date', inplace=True)  
ax = train_data.resample('m')['rating'].count().plot()  
ax.set_title('No of ratings per month (Training data)')  
plt.xlabel('Month')  
plt.ylabel('No of ratings (per month)')  
ax.set_yticklabels([beautify(item, 'M') for item in ax.get_yticks()])
```

```
plt.show()
```




We have sudden increase in number of movies being rated between start of 2004 and end of 2005

```
[112]: no_of Rated movies per user = train_data.groupby(by='user_id')['rating'].  
        ↪count().sort_values(ascending=False)  
  
no_of Rated movies per user.head(10)
```

```
[112]: user_id  
3321    1565  
1333    1087  
3998     805  
5980     753  
4905     689  
9557     639  
2213     624  
4597     622  
6629     618  
8121     595  
Name: rating, dtype: int64
```

```
[113]: fig = plt.figure(figsize=plt.figaspect(.5))
ax1 = plt.subplot(121)
sns.kdeplot(no_of Rated movies_per_user, shade=True, ax=ax1)
plt.xlabel('No of ratings by user')
plt.title("PDF")
```

```
[113]: Text(0.5, 1.0, 'PDF')
```



From the above PDF we can see that most number of users rated less than 1000 number of movies and very few users rated more than 1250 movies.

```
[114]: no_of Rated movies_per_user.describe()
```

```
[114]: count    1803.000000
mean       72.535774
std        105.776838
min         1.000000
25%        13.000000
50%        36.000000
75%        88.500000
```

```
max      1565.000000
Name: rating, dtype: float64
```

From the above summary statistics we can see that 75% of the users rated about 245 movies. That's the reason above PDF is skewed.

```
[115]: np.average(train_data.rating)
```

```
[115]: 3.618043767490939
```

The average rating of all movies is 3.58

4 Creating Sparse Matrix

```
[116]: train_sparse = sparse.csr_matrix((train_data.rating.values, (train_data.user_id.
    ↪values,train_data.movie_id.values)),)
non_zero_tiles = train_sparse.count_nonzero()
a,b = train_sparse.shape
train_sparsity = (1-(non_zero_tiles/(a*b)))*100
```

```
[117]: print("Sparsity of train matrix : ",train_sparsity)
```

```
Sparsity of train matrix : 99.79523720056365
```

```
[118]: test_sparse = sparse.csr_matrix((test_data.rating.values, (test_data.user_id.
    ↪values,test_data.movie_id.values)),)
non_zero_tiles = test_sparse.count_nonzero()
a,b = test_sparse.shape
test_sparsity = (1-(non_zero_tiles/(a*b)))*100
```

```
[119]: print("Sparsity of test matrix : ",test_sparsity)
```

```
Sparsity of test matrix : 99.94763630424391
```

5 Cold Start Problem

```
[120]: total_users = len(np.unique(full_data.user_id))
users_train = len(np.unique(train_data.user_id))
new_users = total_users - users_train
print('Total number of Users :', total_users)
print('Number of Users in Train data :', users_train)
print("No of Users that didn't appear in train data: {}({} %) \n ".
    ↪format(new_users,np.round((new_users/total_users)*100, 2)))
```

```
Total number of Users : 476041
Number of Users in Train data : 1803
```

No of Users that didn't appear in train data: 474238(99.62 %)

```
[121]: total_movies = len(np.unique(full_data.movie_id))
movies_train = len(np.unique(train_data.movie_id))
new_movies = total_movies - movies_train
print('Total number of Movies  :', total_movies)
print('Number of movies in Train data :', movies_train)
print("No of movies that didn't appear in train data: {}({} %) \n ".
      ↪format(new_movies,np.round((new_movies/total_movies)*100, 2)))
```

Total number of Movies : 6386

Number of movies in Train data : 4871

No of movies that didn't appear in train data: 1515(23.72 %)

6 Similarity techniques

6.0.1 We are not going to compute user-user similarity matrix as we have 400k users with each user having 17k dimensional vector it takes more time

6.0.2 Instead we are going to compute the movie-movie similarity matrix which is a (17kx17k) matrix

```
[122]: movie_movie_similarity_matrix = cosine_similarity(train_sparse.T,
      ↪dense_output=False)
movie_movie_similarity_matrix.setdiag(0)
```

```
[123]: movie_movie_similarity_matrix.shape
```

```
[123]: (6387, 6387)
```

We will now store the top 10 similar movies for each of the 17k movies in a dictionary

```
[124]: similar_movies = {}
all_movies = train_data['movie_id'].unique()
for movie_id in all_movies:
    similar_movies_for_movie_id = movie_movie_similarity_matrix[movie_id].
    ↪toarray().ravel().argsort()[::-1][1:][:10]
    similar_movies[movie_id] = similar_movies_for_movie_id
```

```
[125]: movie_titles_master = []
with open('netflix/movie_titles.csv','r') as file:
    for line in file:
        line = line.strip()
        movie_titles_master.append(line)
```



```
[127]: print("Movie details:")
print(movie_titles_master[294])
print("similar movies to 17692 are:")
for movie_id in similar_movies[295]:
    print(movie_titles_master[movie_id-1])
```

Movie details:
295,1995,Ace Ventura: When Nature Calls
similar movies to 17692 are:
2095,1997,Liar Liar
2470,1992,Wayne's World
4393,1994,The Mask: Special Edition
6337,2001,Joe Dirt
5472,1987,Spaceballs
3085,2000,Little Nicky
1509,2002,National Lampoon's Van Wilder
4661,1999,Deuce Bigalow: Male Gigolo
2751,1991,Naked Gun 2 1/2: The Smell of Fear
5318,1995,Tommy Boy

```
[128]: print("Movie details:")
print(movie_titles_master[4392])
print("similar movies to 17627 are:")
for movie_id in similar_movies[4393]:
    print(movie_titles_master[movie_id-1])
```

Movie details:
4393,1994,The Mask: Special Edition
similar movies to 17627 are:
607,1994,Speed
2095,1997,Liar Liar
6350,1998,Rush Hour
2470,1992,Wayne's World
705,1989,Major League
3198,1991,The Addams Family
295,1995,Ace Ventura: When Nature Calls
6196,1984,The Terminator
3648,1988,Who Framed Roger Rabbit?: Special Edition
1659,1993,Grumpy Old Men

```
[129]: print("Movie details:")
print(movie_titles_master[691])
print("similar movies to 9628 are:")
for movie_id in similar_movies[692]:
    print(movie_titles_master[movie_id-1])
```

Movie details:
692,1992,The Hand that Rocks the Cradle

```

similar movies to 9628 are:
3782,1990,Flatliners
607,1994,Speed
4330,1995,While You Were Sleeping
2395,1996,Scream
2095,1997,Liar Liar
4256,1984,Footloose: Special Collector's Edition
5628,1996,The Nutty Professor
2594,1989,Look Who's Talking
4705,1987,Overboard
2462,1987,Planes, Trains and Automobiles

```

```

[130]: print("Movie details:")
print(movie_titles_master[2121])
print("similar movies to 2395 are:")
for movie_id in similar_movies[2122]:
    print(movie_titles_master[movie_id-1])

```

```

Movie details:
2122,1999,Being John Malkovich
similar movies to 2395 are:
571,1999,American Beauty
5226,1998,Rushmore
5862,2000,Memento
5926,1999,Fight Club
5614,2000,Best in Show
6029,2001,Amelie
1865,2004,Eternal Sunshine of the Spotless Mind
788,1994,Clerks
175,1992,Reservoir Dogs
6099,1979,Apocalypse Now

```

We can see that this is a basic algorithm for similarity technique and works good. We will try to build more advanced recommender systems in the upcoming code

7 MACHINE LEARNING MODELS

Building train and test sets in Surprise format for surprise models

```

[131]: from surprise import Reader, Dataset, BaselineOnly

```

```

[132]: data_reader = Reader(rating_scale=(1,5))
train_dataset = Dataset.load_from_df(train_data[['user_id', 'movie_id',
↪ 'rating']], data_reader)
surprise_train_set = train_dataset.build_full_trainset()
surprise_test_set = list(zip(test_data.user_id.values, test_data.movie_id.
↪ values, test_data.rating.values))

```

```
[133]: def get_errors(predictions, print_them=False):

    actual, pred = get_ratings(predictions)
    rmse = np.sqrt(np.mean((pred - actual)**2))
    mape = np.mean(np.abs(pred - actual)/actual)

    return rmse, mape*100

def get_ratings(predictions):
    actual = np.array([pred.r_ui for pred in predictions])
    pred = np.array([pred.est for pred in predictions])

    return actual, pred
```

```
[134]: d = {}
```

8 SURPRISE BASE LINE MODEL

```
[135]: surprise_baseline = BaselineOnly(bsl_options = {'method': 'sgd',
    ↪ 'learning_rate': .01})
surprise_baseline.fit(surprise_train_set)
train_preds = surprise_baseline.test(surprise_train_set.build_testset())
train_actual_ratings= np.array([pred.r_ui for pred in train_preds])
train_bsl_pred_ratings = np.array([pred.est for pred in train_preds])
train_rmse, train_mape = get_errors(train_preds)
print("Train Performance Metrics")
print("Train RMSE:", train_rmse, "Train MAPE:", train_mape)
test_preds = surprise_baseline.test(surprise_test_set)
test_actual_ratings= np.array([pred.r_ui for pred in test_preds])
surprise_baseline_test_pred_ratings = np.array([pred.est for pred in
    ↪ test_preds])
test_rmse, test_mape = get_errors(test_preds)
print("Test Performance Metrics")
print("Test RMSE:", test_rmse, "Test MAPE:", test_mape)

d['Baseline'] = {'Train RMSE': 0.9194028262291728, 'Train MAPE': 27.
    ↪ 98829708497626 , 'Test RMSE': 0.9968439147788294, 'Test MAPE': 30.
    ↪ 640199173819415}
```

Estimating biases using sgd...

Train Performance Metrics

Train RMSE: 0.9194028262291728 Train MAPE: 27.98829708497626

Test Performance Metrics

Test RMSE: 0.9968439147788294 Test MAPE: 30.640199173819415

9 SURPRISE KNN BASELINE ON USER USER SIMILARITY

```
[136]: from surprise import KNNBaseline

[137]: sim_options = {'user_based' : True,
                    'name': 'pearson_baseline',
                    'shrinkage': 100,
                    'min_support': 2
                    }
bsl_options = {'method': 'sgd'}
surprise_knn_bsl_u = KNNBaseline(k=40, sim_options = sim_options, bsl_options =
    ↪bsl_options)
surprise_knn_bsl_u.fit(surprise_train_set)
train_preds = surprise_knn_bsl_u.test(surprise_train_set.build_testset())
train_actual_ratings= np.array([pred.r_ui for pred in train_preds])
train_knn_u_pred_ratings = np.array([pred.est for pred in train_preds])
train_rmse, train_mape = get_errors(train_preds)
print("Train Performance Metrics")
print("Train RMSE:",train_rmse,"Train MAPE:",train_mape)
test_preds = surprise_knn_bsl_u.test(surprise_test_set)
test_actual_ratings= np.array([pred.r_ui for pred in test_preds])
surprise_knn_bsl_u_test_pred_ratings = np.array([pred.est for pred in
    ↪test_preds])
test_rmse, test_mape = get_errors(test_preds)
print("Test Performance Metrics")
print("Test RMSE:",test_rmse,"Test MAPE:",test_mape)
```

Estimating biases using sgd...

Computing the pearson_baseline similarity matrix...

Done computing similarity matrix.

Train Performance Metrics

Train RMSE: 0.450726817425015 Train MAPE: 12.328700243417542

Test Performance Metrics

Test RMSE: 0.9950219958615862 Test MAPE: 30.62410946807211

10 SURPRISE KNN BASELINE ON MOVIE MOVIE SIMILARITY

```
[138]: sim_options = {'user_based' : False,
                    'name': 'pearson_baseline',
                    'shrinkage': 100,
                    'min_support': 2
                    }
bsl_options = {'method': 'sgd'}
surprise_knn_bsl_m = KNNBaseline(k=40, sim_options = sim_options, bsl_options =
    ↪bsl_options)
```

```

surprise_knn_bsl_m.fit(surprise_train_set)
train_preds = surprise_knn_bsl_m.test(surprise_train_set.build_testset())
train_actual_ratings= np.array([pred.r_ui for pred in train_preds])
train_knn_m_pred_ratings = np.array([pred.est for pred in train_preds])
train_rmse, train_mape = get_errors(train_preds)
print("Train Performance Metrics")
print("Train RMSE:",train_rmse,"Train MAPE:",train_mape)
test_preds = surprise_knn_bsl_m.test(surprise_test_set)
test_actual_ratings= np.array([pred.r_ui for pred in test_preds])
surprise_knn_bsl_m_test_pred_ratings = np.array([pred.est for pred in
↪test_preds])
test_rmse, test_mape = get_errors(test_preds)
print("Test Performance Metrics")
print("Test RMSE:",test_rmse,"Test MAPE:",test_mape)

```

Estimating biases using sgd...

Computing the pearson_baseline similarity matrix...

Done computing similarity matrix.

Train Performance Metrics

Train RMSE: 0.4708185685302442 Train MAPE: 12.96997557759143

Test Performance Metrics

Test RMSE: 0.9950219958615862 Test MAPE: 30.62410946807211

11 MATRIX FACTORIZATION TECHNIQUES

SVD Matrix Factorization

[139]: `from surprise import SVD`

[140]:

```

surprise_svd = SVD(n_factors=100, biased=True, random_state=15, verbose=True)
surprise_svd.fit(surprise_train_set)
train_preds = surprise_svd.test(surprise_train_set.build_testset())
train_actual_ratings= np.array([pred.r_ui for pred in train_preds])
train_svd_pred_ratings = np.array([pred.est for pred in train_preds])
train_rmse, train_mape = get_errors(train_preds)
print("Train Performance Metrics")
print("Train RMSE:",train_rmse,"Train MAPE:",train_mape)
test_preds = surprise_svd.test(surprise_test_set)
test_actual_ratings= np.array([pred.r_ui for pred in test_preds])
surprise_svd_test_pred_ratings = np.array([pred.est for pred in test_preds])
test_rmse, test_mape = get_errors(test_preds)
print("Test Performance Metrics")
print("Test RMSE:",test_rmse,"Test MAPE:",test_mape)

```

Processing epoch 0

Processing epoch 1

Processing epoch 2

```

Processing epoch 3
Processing epoch 4
Processing epoch 5
Processing epoch 6
Processing epoch 7
Processing epoch 8
Processing epoch 9
Processing epoch 10
Processing epoch 11
Processing epoch 12
Processing epoch 13
Processing epoch 14
Processing epoch 15
Processing epoch 16
Processing epoch 17
Processing epoch 18
Processing epoch 19
Train Performance Metrics
Train RMSE: 0.677443473528694 Train MAPE: 20.26660309849283
Test Performance Metrics
Test RMSE: 0.9938580620026572 Test MAPE: 30.505403653784867

SVD Matrix Factorization with implicit feedback from user

```

```
[141]: from surprise import SVDpp
```

```
[142]: surprise_svdpp = SVDpp(n_factors=100, random_state=15, verbose=True)
surprise_svdpp.fit(surprise_train_set)
train_preds = surprise_svdpp.test(surprise_train_set.build_testset())
train_actual_ratings= np.array([pred.r_ui for pred in train_preds])
train_svdpp_pred_ratings = np.array([pred.est for pred in train_preds])
train_rmse, train_mape = get_errors(train_preds)
print("Train Performance Metrics")
print("Train RMSE:",train_rmse,"Train MAPE:",train_mape)
test_preds = surprise_svdpp.test(surprise_test_set)
test_actual_ratings= np.array([pred.r_ui for pred in test_preds])
surprise_svdpp_test_pred_ratings = np.array([pred.est for pred in test_preds])
test_rmse, test_mape = get_errors(test_preds)
print("Test Performance Metrics")
print("Test RMSE:",test_rmse,"Test MAPE:",test_mape)

```

```

processing epoch 0
processing epoch 1
processing epoch 2
processing epoch 3
processing epoch 4
processing epoch 5
processing epoch 6

```

```

processing epoch 7
processing epoch 8
processing epoch 9
processing epoch 10
processing epoch 11
processing epoch 12
processing epoch 13
processing epoch 14
processing epoch 15
processing epoch 16
processing epoch 17
processing epoch 18
processing epoch 19
Train Performance Metrics
Train RMSE: 0.5441391507048091 Train MAPE: 15.72906834952557
Test Performance Metrics
Test RMSE: 0.9942700695822917 Test MAPE: 30.490688418774592

```

12 ENSEMBLE MODEL

```

[143]: all_test_predictions = np.
        ↪array([surprise_baseline_test_pred_ratings,surprise_knn_bsl_u_test_pred_ratings,surprise_kn
all_test_preds = np.mean(all_test_predictions,axis=0)
test_actual_ratings

```

```

[143]: array([3, 3, 3, ..., 5, 5, 4], dtype=int64)

```

```

[144]: import numpy as np

def rmse(y_true, y_pred):
    return np.sqrt(np.mean((y_true - y_pred) ** 2))

def mape(y_true, y_pred):
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

test_rmse = rmse(test_actual_ratings, all_test_preds)
test_mape = mape(test_actual_ratings, all_test_preds)

print("Test RMSE:", test_rmse)
print("Test MAPE:", test_mape)

```

```

Test RMSE: 0.9939049783542391
Test MAPE: 30.543984512349482

```

13 FEATURIZING THE DATA FOR REGRESSION MODELS

```
[145]: test_data['bsl'] = surprise_baseline_test_pred_ratings
test_data['knn_u'] = surprise_knn_bsl_u_test_pred_ratings
test_data['knn_m'] = surprise_knn_bsl_m_test_pred_ratings
test_data['svd'] = surprise_svd_test_pred_ratings
test_data['svdpp'] = surprise_svdpp_test_pred_ratings
```

```
[146]: test_data.shape
```

```
[146]: (43448, 9)
```

```
[147]: test_data.head()
```

```
[147]:
```

	movie_id	user_id	rating	date	bsl	knn_u	knn_m	\
	13803346	2640	10268	3	2000-01-06	4.079549	4.036668	4.036668
	5894733	1157	10268	3	2000-01-06	3.362539	3.427967	3.427967
	4698159	931	10268	3	2000-01-06	3.280227	3.402718	3.402718
	29713806	5444	10268	2	2000-01-06	3.441228	3.518444	3.518444
	35113306	6285	10268	3	2000-01-06	3.358916	3.446778	3.446778

	svd	svdpp
13803346	3.932473	3.804996
5894733	3.448295	3.422181
4698159	3.435544	3.405568
29713806	3.539112	3.516245
35113306	3.454049	3.437214

```
[148]: train_data['bsl'] = train_bsl_pred_ratings
train_data['knn_u'] = train_knn_u_pred_ratings
train_data['knn_m'] = train_knn_m_pred_ratings
train_data['svd'] = train_svd_pred_ratings
train_data['svdpp'] = train_svdpp_pred_ratings
```

```
[149]: train_data.shape
```

```
[149]: (130782, 8)
```

```
[150]: train_data.head()
```

```
[150]:
```

	movie_id	user_id	rating	bsl	knn_u	knn_m	svd	\
date								
1999-12-31	295	1086	4	3.186765	3.365851	3.348955	2.959398	
1999-12-31	4652	1086	3	3.097945	3.155964	3.281321	3.071056	
1999-12-31	829	1086	3	3.157202	3.229965	3.283076	3.019874	
1999-12-31	5237	1086	5	3.237026	3.907288	4.018165	3.694763	
1999-12-31	682	1086	4	3.865532	3.924292	3.955110	3.907480	


```

            svdpp
date
1999-12-31  3.147548
1999-12-31  3.122768
1999-12-31  3.084087
1999-12-31  4.026036
1999-12-31  4.029154

```

```

[151]: global_user_avg = test_data.groupby('user_id')['rating'].mean()
global_movie_avg = test_data.groupby('movie_id')['rating'].mean()
test_data['global_avg_user_rating'] = test_data['user_id'].map(global_user_avg)
test_data['global_avg_movie_rating'] = test_data['movie_id'].
    ↪map(global_movie_avg)

```

```

[152]: global_user_avg = train_data.groupby('user_id')['rating'].mean()
global_movie_avg = train_data.groupby('movie_id')['rating'].mean()
train_data['global_avg_user_rating'] = train_data['user_id'].
    ↪map(global_user_avg)
train_data['global_avg_movie_rating'] = train_data['movie_id'].
    ↪map(global_movie_avg)

```

```

[153]: train_data.head()

```

```

[153]:
            movie_id  user_id  rating      bsl      knn_u      knn_m      svd  \
date
1999-12-31         295     1086        4  3.186765  3.365851  3.348955  2.959398
1999-12-31        4652     1086        3  3.097945  3.155964  3.281321  3.071056
1999-12-31         829     1086        3  3.157202  3.229965  3.283076  3.019874
1999-12-31        5237     1086        5  3.237026  3.907288  4.018165  3.694763
1999-12-31         682     1086        4  3.865532  3.924292  3.955110  3.907480

```

```

            svdpp  global_avg_user_rating  global_avg_movie_rating
date
1999-12-31  3.147548                3.56                3.447619
1999-12-31  3.122768                3.56                3.083333
1999-12-31  3.084087                3.56                3.166667
1999-12-31  4.026036                3.56                3.295918
1999-12-31  4.029154                3.56                3.846154

```

```

[154]: test_data.head()

```

```

[154]:
            movie_id  user_id  rating      date      bsl      knn_u      knn_m  \
13803346         2640     10268        3  2000-01-06  4.079549  4.036668  4.036668
5894733          1157     10268        3  2000-01-06  3.362539  3.427967  3.427967
4698159          931     10268        3  2000-01-06  3.280227  3.402718  3.402718
29713806         5444     10268        2  2000-01-06  3.441228  3.518444  3.518444

```

35113306	6285	10268	3	2000-01-06	3.358916	3.446778	3.446778
	svd	svdpp		global_avg_user_rating		global_avg_movie_rating	
13803346	3.932473	3.804996		3.177305		3.222222	
5894733	3.448295	3.422181		3.177305		2.000000	
4698159	3.435544	3.405568		3.177305		3.000000	
29713806	3.539112	3.516245		3.177305		2.666667	
35113306	3.454049	3.437214		3.177305		3.000000	

```
[155]: import xgboost as xgb
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc, mean_squared_error
from sklearn.model_selection import GridSearchCV
```

```
[156]: y_train = train_data['rating']
X_train = train_data.drop(columns=['rating'])
y_test = test_data['rating']
X_test = test_data.drop(columns=['rating'])
```

```
[157]: X_test.head()
```

```
[157]:
```

	movie_id	user_id	date	bsl	knn_u	knn_m	\
	13803346	2640	10268	2000-01-06	4.079549	4.036668	4.036668
	5894733	1157	10268	2000-01-06	3.362539	3.427967	3.427967
	4698159	931	10268	2000-01-06	3.280227	3.402718	3.402718
	29713806	5444	10268	2000-01-06	3.441228	3.518444	3.518444
	35113306	6285	10268	2000-01-06	3.358916	3.446778	3.446778

	svd	svdpp	global_avg_user_rating	global_avg_movie_rating
13803346	3.932473	3.804996	3.177305	3.222222
5894733	3.448295	3.422181	3.177305	2.000000
4698159	3.435544	3.405568	3.177305	3.000000
29713806	3.539112	3.516245	3.177305	2.666667
35113306	3.454049	3.437214	3.177305	3.000000

14 XGBoost

```
[158]: max_depth = np.array([3, 6, 9, 12, 15])
learning_rate = np.array([0.01, 0.05, 0.1])
n_estimators = np.array([20, 35, 50, 70, 100])
params_dict = [{'max_depth': max_depth, 'learning_rate': learning_rate,
'n_estimators': n_estimators}]
xgb_optimal = xgb.XGBRegressor(random_state=15, n_jobs=-1)
grid = GridSearchCV(estimator=xgb_optimal,
                    param_grid=params_dict,
                    scoring='neg_mean_squared_error', n_jobs=-1, cv=5,
```

```

        return_train_score=True)
grid_result = grid.fit(X_train, y_train)
train_score = grid_result.cv_results_['mean_train_score']
train_score_std = grid_result.cv_results_['std_train_score']
cv_score = grid_result.cv_results_['mean_test_score']
cv_score_std = grid_result.cv_results_['std_test_score']
print("Optimal Parameters : ", grid_result.best_estimator_.get_params())
train_score = train_score.reshape(len(learning_rate),len(max_depth),
                                  len(n_estimators))
cv_score = cv_score.reshape(len(learning_rate),len(max_depth),
                             len(n_estimators))
for lr in range(len(learning_rate)):
    plt.figure(figsize=(10.0, 8.0))
    ax = sns.heatmap(train_score[lr,:,:],
                      annot=True, square=False, cmap="Oranges",
                      xticklabels=["nE = "+str(ne) for ne in n_estimators],
                      yticklabels=["mD = "+str(md) for md in max_depth])


    plt.title("Training scores for max_depth and num_estimators "+
              "with learning_rate="+str(learning_rate[lr]))
    plt.show()
    print('')
    plt.figure(figsize=(10.0, 8.0))
    ax = sns.heatmap(cv_score[lr,:,:],
                      annot=True, square=False, cmap="Blues",
                      xticklabels=["nE = "+str(ne) for ne in n_estimators],
                      yticklabels=["md = "+str(md) for md in max_depth])
    plt.title("Cross-val scores for max_depth and num_estimators "+
              "with learning_rate="+str(learning_rate[lr]))
    plt.show()


```


```


Optimal Parameters : {'objective': 'reg:squarederror', 'base_score': None,
'booster': None, 'callbacks': None, 'colsample_bylevel': None,
'colsample_bynode': None, 'colsample_bytree': None, 'device': None,
'early_stopping_rounds': None, 'enable_categorical': False, 'eval_metric': None,
'feature_types': None, 'gamma': None, 'grow_policy': None, 'importance_type':
None, 'interaction_constraints': None, 'learning_rate': 0.05, 'max_bin': None,
'max_cat_threshold': None, 'max_cat_to_onehot': None, 'max_delta_step': None,
'max_depth': 6, 'max_leaves': None, 'min_child_weight': None, 'missing': nan,
'monotone_constraints': None, 'multi_strategy': None, 'n_estimators': 70,
'n_jobs': -1, 'num_parallel_tree': None, 'random_state': 15, 'reg_alpha': None,
'reg_lambda': None, 'sampling_method': None, 'scale_pos_weight': None,
'subsample': None, 'tree_method': None, 'validate_parameters': None,
'verbosity': None}


```

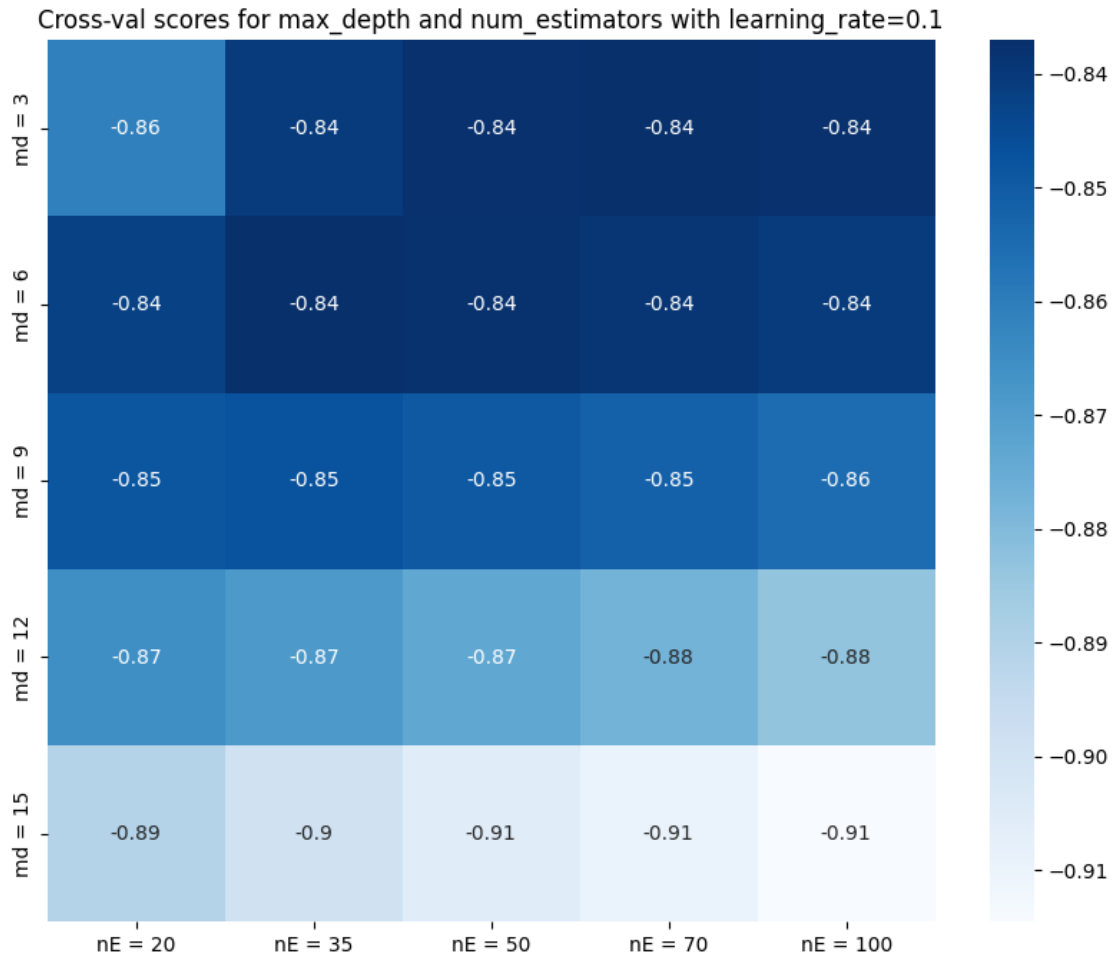












```
[159]: def get_error_metrics(y_true, y_pred):
        rmse = np.sqrt(np.mean([ (y_true[i] - y_pred[i])**2 for i in
        ↪range(len(y_pred)) ]))
        mape = np.mean(np.abs( (y_true - y_pred)/y_true )) * 100
        return rmse, mape
```

```
[160]: X_test = X_test.drop(['date'],axis=1)
```

```
[161]: xgb_optimal = xgb.XGBRegressor(max_depth=6,
                                     learning_rate=0.05,
                                     n_estimators=70,
                                     n_jobs=-1,random_state=15)

xgb_optimal.fit(X_train, y_train)
pred_train = xgb_optimal.predict(X_train)
pred_test = xgb_optimal.predict(X_test)
print("Using max depth value for tree - ",6)
```

```

print("Using num estimators for tree - ", 70)
print("Using learning rate for tree - ", 0.05)

train_rmse, train_mape = get_error_metrics(y_train, pred_train)
print("Train Error Metrics. \n\tRMSE: {}\n\tMAPE: {}".format(
    train_rmse, train_mape))
test_rmse, test_mape = get_error_metrics(np.array(y_test), np.array(pred_test))
print("Test Error Metrics. \n\tRMSE: {}\n\tMAPE: {}".format(
    test_rmse, test_mape))

```


```

Using max depth value for tree - 6
Using num estimators for tree - 70
Using learning rate for tree - 0.05
Train Error Metrics.
    RMSE: 0.9032830684516676
    MAPE: 27.527433383281476
Test Error Metrics.
    RMSE: 0.8637764269642914
    MAPE: 25.15808375618946

```

```
[162]: xgb.plot_importance(xgb_optimal)
```

```
[162]: <Axes: title={'center': 'Feature importance'}, xlabel='F score',
ylabel='Features'>
```



```
[163]: d = {}
d['Baseline'] = {'Train RMSE': 0.9194028262291728, 'Train MAPE': 27.98829708497626, 'Test RMSE': 0.9968439147788294, 'Test MAPE': 30.640199173819415}
d['KNN u-u'] = {'Train RMSE': 0.450726817425015, 'Train MAPE': 12.328700243417542, 'Test RMSE': 0.9950219958615862, 'Test MAPE': 30.62410946807211}
d['knn m-m'] = {'Train RMSE': 0.4708185685302442, 'Train MAPE': 12.96997557759143, 'Test RMSE': 0.9950219958615862, 'Test MAPE': 30.62410946807211}
d['svd'] = {'Train RMSE': 0.677443473528694, 'Train MAPE': 20.26660309849283, 'Test RMSE': 0.9938580620026572, 'Test MAPE': 30.505403653784867}
d['svdpp'] = {'Train RMSE': 0.5441391507048091, 'Train MAPE': 15.72906834952557, 'Test RMSE': 0.9942700695822917, 'Test MAPE': 30.490688418774592}
d['xgb+surprise'] = {'Train RMSE': 0.9032830684516676, 'Train MAPE': 27.527433383281476, 'Test RMSE': 0.8637764269642914, 'Test MAPE': 25.15808375618946}
```

```
[164]: import matplotlib.pyplot as plt

# Define the data
models = ['Baseline', 'KNN u-u', 'KNN m-m', 'SVD', 'SVD++', 'XGB+Surprise']
train_rmse = [0.9194028262291728, 0.450726817425015, 0.4708185685302442, 0.677443473528694, 0.5441391507048091, 0.9032830684516676]
train_mape = [27.98829708497626, 12.328700243417542, 12.96997557759143, 20.26660309849283, 15.72906834952557, 27.527433383281476]
test_rmse = [0.9968439147788294, 0.9950219958615862, 0.9950219958615862, 0.9938580620026572, 0.9942700695822917, 0.8637764269642914]
test_mape = [30.640199173819415, 30.62410946807211, 30.62410946807211, 30.505403653784867, 30.490688418774592, 25.15808375618946]

# Create subplots
fig, axs = plt.subplots(2, 2, figsize=(14, 10))

# Train RMSE
axs[0, 0].bar(models, train_rmse, color='skyblue')
axs[0, 0].set_title('Train RMSE')
axs[0, 0].set_ylabel('RMSE')
for i, v in enumerate(train_rmse):
    axs[0, 0].text(i, v + 0.01, f'{v:.6f}', ha='center')

# Train MAPE
axs[0, 1].bar(models, train_mape, color='lightgreen')
axs[0, 1].set_title('Train MAPE')
axs[0, 1].set_ylabel('MAPE')
for i, v in enumerate(train_mape):
    axs[0, 1].text(i, v + 0.5, f'{v:.6f}', ha='center')
```

```

# Test RMSE
axs[1, 0].bar(models, test_rmse, color='salmon')
axs[1, 0].set_title('Test RMSE')
axs[1, 0].set_ylabel('RMSE')
for i, v in enumerate(test_rmse):
    axs[1, 0].text(i, v + 0.01, f'{v:.6f}', ha='center')

# Test MAPE
axs[1, 1].bar(models, test_mape, color='gold')
axs[1, 1].set_title('Test MAPE')
axs[1, 1].set_ylabel('MAPE')
for i, v in enumerate(test_mape):
    axs[1, 1].text(i, v + 0.5, f'{v:.6f}', ha='center')

# Rotate x-axis labels for better readability
for ax in axs.flat:
    ax.set_xticklabels(models, rotation=45, ha='right')

# Adjust layout
plt.tight_layout()

# Show plot
plt.show()

```

